Design of high fidelity combinatorial libraries

November 4, 2026

Abstract

Recent tools for de novo design have revolutionized the way we propose novel proteins and the types of challenges we are able to tackle. While these tools have dramatically increased the hit rate of functional designs, zero-shot generations often require further optimization to enhance the desired properties. We present a novel method for generating high fidelity diverse combinatorial libraries. Traditional library generation methods rely on error prone PCR or natural fragment shuffling to generate diversity. Such approaches do not take into account the filtering of designs through structure prediction models or other desired oracles. We demonstrate the ability to post-train inverse folding models to design fragments with filter pass rates of 93.4% and libraries which can easily scale beyond 10^{12} individually addressable variants. Our post-training framework enables control over the search space by centering the library around a reference sequence with tunable mutation distance. Once designed, we demonstrate how individual variants are assembled with simple molecular biology techniques (golden gate assembly). Our cell-free assay platform enables us to screen enzymes at the scale of 10^3 measurements per day, each with replicates. The libraries generated are easily much larger than most traditional and modern high throughput assays are able to screen, so we demonstrate how the library can be sub sampled smartly to optimize enzymes. For a de novo designed PETase we demonstrate our platform finds a variant with 250x increase in catalytic activity.

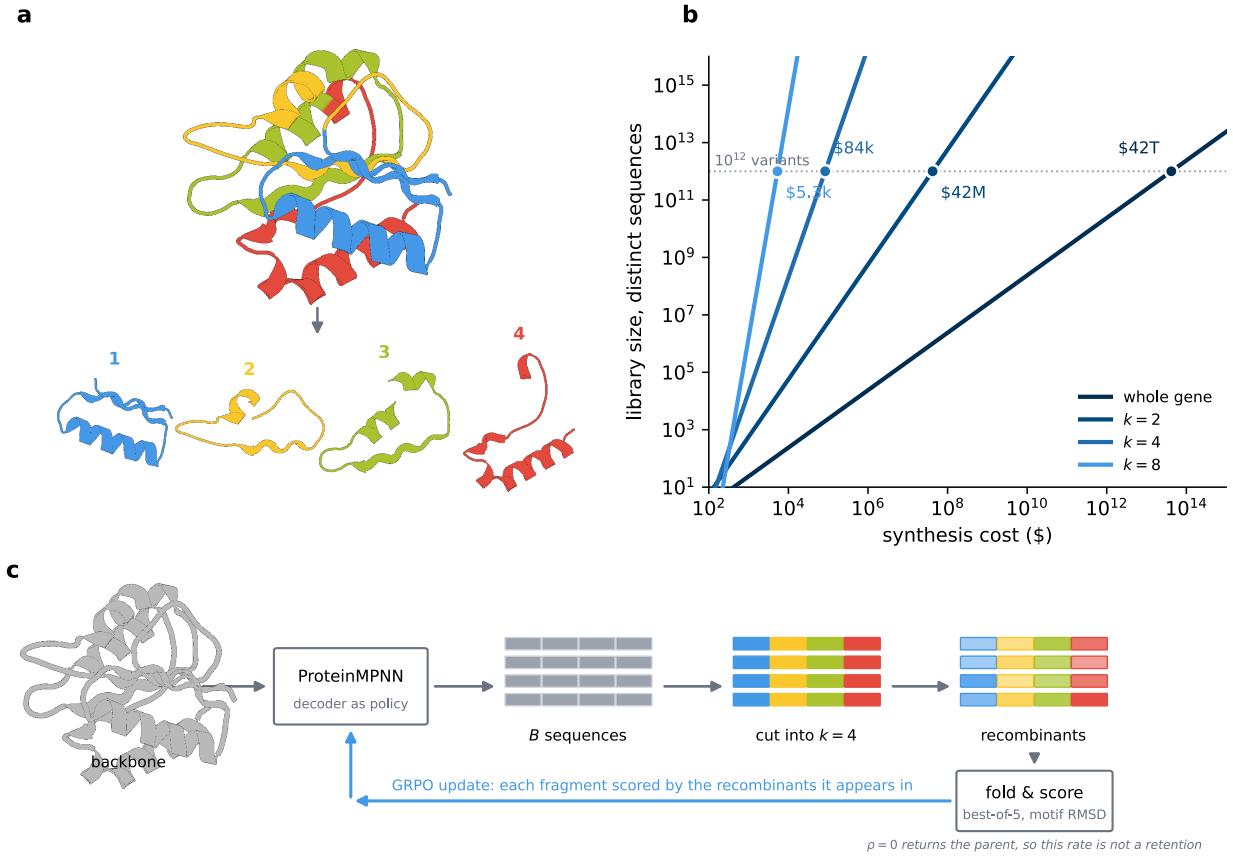


Figure 1: **A combinatorial library, and the policy that makes one assemblable. (a)** One fold cut into $k = 4$ equal fragments, shown on flavodoxin (2FX2, chain A, 147 residues; 39 % helix, 33 % strand). The assembled fold is shown above and its four fragments below, numbered along the chain. All five views share one camera and one scale, so the pieces are the fold's own and their relative sizes are real. Equal lengths are chosen so that assembly is robust, not because the boundaries are structurally meaningful, and at this length the four fragments are 37 residues each — inside the 20 to 50 band the method recommends. **(b)** What a given spend buys. Whole genes cost $42 each at eBlock pricing for a 200-aa enzyme, so library size is linear in spend; a combinatorial library buys k sets of parts instead of N products, and size goes as the kth power of what each slot can hold. eBlocks are billed at a 300-bp minimum, which binds from $k = 2$ on, so past that point cutting finer stops buying shorter DNA and only buys more pieces of it. Reaching 10^{12} variants costs $42T as whole genes, $42M at $k = 2$, $84k at $k = 4$ and $5.3k at $k = 8$. **(c)** The method. A backbone conditions the ProteinMPNN decoder, which is treated as a policy; the sequences it emits are cut at the fragment boundaries, recombined across parents, and folded best-of-5. Each fragment is scored by the recombinants it appears in, and that per-fragment reward drives the GRPO update. The policy is therefore optimised for how well its fragments *recombine*, not for how good any single sequence is.